

Lec 1. From Generation to Sampling.

Formalize "generating"

objects: vectors $\mathbf{z} \in \mathbb{R}^d$. - we generate.

How good a image is \approx How likely it is under data distribution.

$p_{\text{data}}: \mathbb{R}^d \rightarrow \mathbb{R}$. Note: we do NOT know the prob density!

$$\mathbf{z} \mapsto p_{\text{data}}(\mathbf{z})$$

• Generation means sampling from the distribution.

• A dataset consists of finite samples of the distribution.

$$\mathbf{z}_1, \dots, \mathbf{z}_n \sim p_{\text{data}}.$$

• Conditional / Non-conditional generation.

↓
Fixed
prompt.

↓
sampling the conditional data dist.
 $\mathbf{z} \sim p_{\text{data}}(\cdot | y)$

Flexible prompt.

Generative Models generate samples from data distribution



Initial distribution:

$$p_{\text{init}}$$

Default:

$$p_{\text{init}} = \mathcal{N}(0, I_d)$$

A generative model converts samples from a initial distribution (e.g. Gaussian) into samples from the data distribution:

$$\mathbf{x} \sim p_{\text{init}}$$



Generative
Model

$$\mathbf{z} \sim p_{\text{data}}$$



think about p in it as a gorion so as a



How this achieved. This is what we cares about.

Flow Models

Def. Trajectory. $x: [0, 1] \rightarrow \mathbb{R}^d$ $t \mapsto x_t$.



Vector field $u: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$. $(x, t) \mapsto u_t(x)$



ODE: $x_0 = x_0$ (initial condition)

$$\frac{dx}{dt} = u_t(x_0)$$

Flow: A collection of solutions to ODE. (i.e., a collection of trajectory)

$\varphi: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$ $(x_0, t) \mapsto \varphi_t(x_0)$

↑ ↑
Init time.

$$\begin{cases} \varphi_0(x_0) = x_0 \\ \frac{d}{dt} \varphi_t(x_0) = u_t(\varphi_t(x_0)). \end{cases}$$

理解: ODE 由向量场确定.

若我们知道 flow 后, ODE 的解就可

↓
给出了 $\varphi_t(x_0)$,
即 x_t 的位置
↑
Fix init. find trajectory.

Vector fields define ODEs whose solutions are flows.

Existence and Uniqueness Theorem ODEs



Theorem (Picard–Lindelöf theorem): If the vector field $u_t(x)$ is continuously differentiable with bounded derivatives, then a unique solution to the ODE

$$X_0 = x_0, \quad \frac{d}{dt} X_t = u_t(X_t)$$

↳ Lipschitz.

exists. In other words, a flow map exists. More generally, this is true if the vector field is **Lipschitz**.

Key takeaway: In the cases of practical interest for machine learning, **unique solutions to ODE/flows exist.**

Eg. Linear ODE.

$$u_t(x) = -\theta x \quad (\theta > 0)$$

Flow given by: $\psi_t(x_0) = \exp(-\theta t) x_0$.

Check in it
vector field.

Pf: (Verification)

$$\psi_0(x_0) = \exp(0)x_0 = x_0 \quad \checkmark$$

$$\frac{d}{dt} \psi_t(x_0) = -\theta x_0 \exp(-\theta t) = -\theta \psi_t(x_0) = u_t(\psi_t(x_0)) \quad \checkmark$$

Numerical ODE simulation - Euler method



Algorithm 1 Simulating an ODE with the Euler method

Require: Vector field u_t , initial condition x_0 , number of steps n

- 1: Set $t = 0$
- 2: Set step size $h = \frac{1}{n}$
- 3: Set $X_0 = x_0$
- 4: **for** $i = 1, \dots, n - 1$ **do**
- 5: $X_{t+h} = X_t + h u_t(X_t)$ *Small step into direction of vector field*
- 6: Update $t \leftarrow t + h$
- 7: **end for**
- 8: **return** $X_0, X_h, X_{2h}, \dots, X_1$ *Return trajectory*

Flow model.

$$p_{\text{init}} \xrightarrow{\text{ODE}} p_{\text{data}}$$

Neural network: $u_t^\theta : \mathbb{R}^d \times [0,1] \rightarrow \mathbb{R}^d$.

θ : parameter,

befk ... to specify vector field nn.

Random init: $X_0 \sim p_{\text{init}}$.

$$\text{ODE: } \frac{d}{dt} X_t = u_t^\theta(X_t)$$

Goal: $X_1 \sim p_{\text{data}}$
(end)

How to generate objects with a Flow Model



Algorithm 1 Sampling from a Flow Model with Euler method

Require: Neural network vector field u_t^θ , number of steps n

- 1: Set $t = 0$
- 2: Set step size $h = \frac{1}{n}$
- 3: Draw a sample $X_0 \sim p_{\text{init}}$ *Random initialization!*
- 4: **for** $i = 1, \dots, n - 1$ **do**
- 5: $X_{t+h} = X_t + h u_t^\theta(X_t)$
- 6: Update $t \leftarrow t + h$
- 7: **end for**
- 8: **return** X_1 *Return final point*

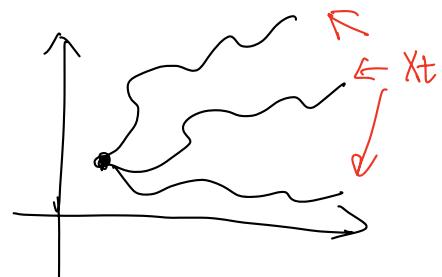
2.2. Diffusion.

SDE

The solution is random trajectory. (Can be seen as stochastic process).

X_t : random var. (即使 x_0, t 确定, X_t 仍不确定).

$$X: [0,1] \rightarrow \mathbb{R}^d. t \mapsto X_t$$



Vector field.

$$u \in \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d. + \underbrace{\text{Diffusion coefficient}}$$

$$\sigma: [0, 1] \rightarrow \mathbb{R}_{\geq 0}$$

put randomness

into SDE.

SDE:

$$X_0 = x_0 \quad \underline{\text{init cond.}} \quad (\text{We fix } x_0).$$

$$dX_t = u_t(X_t) dt + \sigma_t dW_t$$

We: Brownian motion.

Brownian motion. (A specific Gaussian process)

$$W = (W_t)_{t \geq 0} \in \mathbb{R}^d.$$

$$\mathbb{E} W_0 = 0.$$

$$\textcircled{2} \text{ Gaussian increments} \quad W_t - W_s \sim N(0, (t-s) Id) \quad 0 \leq s < t.$$

\textcircled{3} Independent increments.

$$W_{t_1} - W_{t_0}, \dots, W_{t_n} - W_{t_{n-1}} \quad 0 \leq t_0 < t_1 < \dots < t_n.$$

are independent.

↑
linear increased Variance.

problems: Brownian motion has no derivatives.

" dX_t " notation:

$$\textcircled{1} \quad \frac{d}{dt} X_t = u_t(X_t) \Leftrightarrow X_{t+h} = X_t + h u_t(X_t) + \underbrace{h \sigma_t(h)}_{O(h)}$$

$$\boxed{\lim_{h \rightarrow 0} \sigma_t(h) \rightarrow 0}$$

Think this as a Taylor approximation.

$$\text{pf:} \quad \lim_{h \rightarrow 0} \frac{X_{t+h} - X_t}{h} = u_t(X_t) \Rightarrow \frac{X_{t+h} - X_t}{h} = u_t(X_t) + \sigma_t(h). \quad \square$$

$$2. \quad dX_t = u_t(X_t) dt + \sigma_t dW_t \Leftrightarrow X_{t+h} = X_t + h u_t(X_t) + \sigma_t (W_{t+h} - W_t) + h R_t(h)$$

$W_{t+h} - W_t \sim \mathcal{N}(0, h I_d)$

Sampling.

$\left(\lim_{h \rightarrow 0} \sqrt{\mathbb{E}[|R_t(h)|^2]} = 0, \right)$
 $(\text{Var}(R_t(h)) = 0)$

Existence and Uniqueness Theorem SDEs



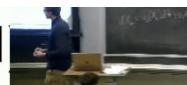
Theorem: If the vector field $u_t(x)$ is continuously differentiable with bounded derivatives and the diffusion coeff. is continuous, then a *unique* solution to the SDE

$$X_0 = x_0, \quad dX_t = u_t(X_t)dt + \sigma_t dW_t$$

exists. More generally, this is true if the vector field is **Lipschitz**.

Key takeaway: In the cases of practical interest for machine learning, **unique solutions to SDE**.

Numerical SDE simulation (Euler-Maruyama method)



Algorithm 2 Sampling from a SDE (Euler-Maruyama method)

Require: Vector field u_t , number of steps n , diffusion coefficient σ_t

- 1: Set $t = 0$
 - 2: Set step size $h = \frac{1}{n}$
 - 3: Set $X_0 = x_0$
 - 4: **for** $i = 1, \dots, n-1$ **do**
 - 5: Draw a sample $\epsilon \sim \mathcal{N}(0, I_d)$
 - 6: $X_{t+h} = X_t + h u_t(X_t) + \sigma_t \sqrt{h} \epsilon$
 - 7: Update $t \leftarrow t + h$
 - 8: **end for**
 - 9: **return** $X_0, X_h, X_{2h}, X_{3h}, \dots, X_1$
- This is just 1 possible process.

σ_t : Can be dependent on X ? (state-dependent)

Eq. ODE-process

$$dx_t = -\theta x_t dt + \sigma dW_t.$$

Diffusion model:

$$p_{\text{init}} \xrightarrow{\text{SDE}} p_{\text{data}}.$$

$$\text{NN: } u_t^\theta: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d.$$

Diffusion coefficient: σ_t (fixed)

Random init : $x_0 \sim \text{input}$.

Goal: $x_1 \sim p_{\text{data}}$.

$$\text{SDE : } u_t^\theta(x_t) dt + \sigma_t dW_t$$

Algorithm 2 Sampling from a Diffusion Model (Euler-Maruyama method)

Require: Neural network u_t^θ , number of steps n , diffusion coefficient σ_t

- 1: Set $t = 0$
 - 2: Set step size $h = \frac{1}{n}$
 - 3: Draw a sample $X_0 \sim p_{\text{init}}$
 - 4: **for** $i = 1, \dots, n - 1$ **do**
 - 5: Draw a sample $\epsilon \sim \mathcal{N}(0, I_d)$
 - 6: $X_{t+h} = X_t + h u_t^\theta(X_t) + \sigma_t \sqrt{h} \epsilon$
 - 7: Update $t \leftarrow t + h$
 - 8: **end for**
 - 9: **return** X_1
-